

Decision-Optimal Probabilistic Forecast Reconciliation for Expected Shortfall in Hierarchical Time Series

Vincent Su

1 Background

Risk managers and regulators use quantitative risk measures to quantify and manage potential losses in uncertain environments. Two familiar choices are Value-at-Risk (VaR)—a loss threshold exceeded with some probability—and Expected Shortfall (ES; also known as Conditional VaR)—the average loss beyond that threshold.

These tail risks are, however, not managed in isolation. Consider a hedge fund’s investment portfolio returns, which consists returns of different asset classes, and further contains returns of various individual securities. This is an example of a hierarchical structure, with multiple *aggregation levels*. Risk managers need to assess and forecast the risk quantities at several levels, from the overall portfolio, down to individual securities. Ignoring this naturally-formed hierarchical structure can lead to misaligned forecasts, for example, the forecasted returns of individual securities are low while the overall portfolio return is high. This results in misallocation of capital and suboptimal risk management strategies.

Forecasting such hierarchical structure also arises in many other decision-making contexts, from supply chains (Angam et al., 2025; Seaman & Bowman, 2022) and energy planning (Di Modica et al., 2021), to macroeconomic (El Gemayel et al., 2022; Li et al., 2019) and tourism analysis (Athanasopoulos et al., 2009). Stakeholders in these settings need forecasts at several aggregation levels to allocate resources and manage risk. These hierarchical structures naturally form *aggregation constraints* (child levels must sum up to parents). Taken together, when we have a collection of time series organised in a hierarchy subject to aggregation constraints, we refer to this as *hierarchical time series*.

In practice, when forecasts are produced for all series (often called *base forecasts*), they typically violate the aggregation constraints observed in the data; such forecasts are *incoherent* (e.g., security-level numbers don’t sum to the overall trading book). This can undermine downstream decisions that require internal consistency. Forecast reconciliation, a post-processing step, addresses this by adjusting the base forecasts so the final adjusted set of forecasts is *coherent* with the aggregation structure. Approaches such as OLS (Hyndman et al., 2011), WLS (Hyndman et al., 2016), MinT (Wickramasuriya et al., 2019), and more are developed and commonly used in practice. An additional incentive for the rapid adoption of reconciliation methods is that they improve forecast accuracy (Athanasopoulos et al., 2024). For instance, the International Monetary Fund reported improvements in their liquidity forecasts,

generating economic benefits via more efficient monetary policy operations (El Gemayel et al., 2022).

Importantly, ES behaves well under aggregation (satisfying subadditivity¹), making it a natural target when risks are rolled up through a hierarchy. This context motivates methods that deliver coherent, multi-level forecasts whose uncertainty aligns with the tail-risk metrics managers actually use.

(should i remove this?)

2 Motivation

Modern risk management and operations need more than a typical mean or median value—they need quantities from the full predictive distribution that capture uncertainty and tail behaviour, because many decisions are tail-sensitive (e.g., extreme demand in retail settings or energy generation shortfalls). While VaR is conceptually straightforward, it ignores the severity of losses beyond the VaR threshold and is not a coherent risk measure, since it does not satisfy the subadditivity property (Yamai & Yoshida, 2005). VaR can mislead risk assessments across different levels of the hierarchy, and may underestimate the true risk exposure under market stress. ES, on the other hand, is a coherent metric that captures the average loss in the tail beyond the VaR threshold. This makes ES a more reliable metric for assessing and managing “tail risks”, and it has been adopted as the preferred regulatory metric since the Basel III framework (Basel Committee on Banking Supervision, 2019). If ES is the metric that risk managers actually use, reconciliation methods should guarantee the coherence and optimality with respect to ES.

Most reconciliation methodology has been developed for point forecasts and are optimised for point scoring functions such as mean squared error (MSE). Generalising these methods to reconcile probabilistic forecasts does not guarantee optimality for tail-risk measures such as VaR or ES. Although Panagiotelis et al. (2023) introduced an score optimisation for reconciliation weights using proper scoring rules (energy score and variogram score) for full multivariate predictive distributions, they do not directly address the tail functionals. This leaves the question: *“How can we construct probabilistic forecast reconciliation that are explicitly optimised for tail-risk measures such as ES?”*

A practical challenge is that ES is not elicitable on its own, meaning there is no scoring rule or loss function that can uniquely identify ES as the optimiser. However, ES is jointly elicitable with VaR (Fissler et al., 2015), and is amenable to convex optimisation via an auxiliary VaR

¹**Subadditivity.** A risk measure ρ on loss variables is subadditive if $\rho(X + Y) \leq \rho(X) + \rho(Y)$ for any losses X, Y . It encodes diversification: the risk of a combined portfolio should not exceed the sum of stand-alone risks. Expected Shortfall satisfies subadditivity for all $\alpha \in (0, 1)$, whereas VaR can violate it. Thus, the portfolio ES cannot exceed the sums of desk ESs.

variable (Rockafellar & Uryasev, 2000). This opens the door to train reconciliation specifically for ES.

As a result, the objective of this project is to bring these theoretical advances into the reconciliation framework and develop methods that produce coherent probabilistic forecasts optimised for ES. Our goal is a reliable, consistent, and decision-aligned risk-forecasting framework that helps national and international organisations improve ES backtesting, stabilise capital, and allocate resources more effectively across hierarchical levels.

3 Methodology

Preliminaries

Let $\mathbf{y}_t \in \mathbb{R}^K$ denote a vector of K time series variables from the hierarchy, at time t . For simplicity, we suppress the subscript t and denote it as \mathbf{y} . These variables are subject to linear aggregation constraints represented by a matrix \mathbf{C} , such that $\mathbf{C}\mathbf{y} = 0$. For example, in a retail setting, the total sales at a warehouse (y_1) and two stores (y_2 and y_3) might satisfy the constraint $y_1 = y_2 + y_3$, leading to $\mathbf{C} = (1, -1, -1)$.

Base forecasts $\hat{\mathbf{y}}$, which we generate independently for each series, are typically *incoherent* (i.e. $\mathbf{C}\hat{\mathbf{y}} \neq 0$), while reconciled forecasts $\tilde{\mathbf{y}}$ are coherent by construction (i.e. $\mathbf{C}\tilde{\mathbf{y}} = 0$). We obtain coherence via a linear mapping $\tilde{\mathbf{y}} = \mathbf{P}\hat{\mathbf{y}}$, where \mathbf{P} maps the full *incoherent* space \mathbb{R}^K to the linear *coherent* subspace $\mathbb{S} \subset \mathbb{R}^K$, where aggregation constraints hold.

Optimisation Framework

For a random variable Y , the Expected Shortfall (ES) at level $\alpha \in (0, 1)$ is defined as $\text{ES}_\alpha(Y) = E[Y \mid Y < \text{VaR}_\alpha(Y)]$, where $\text{VaR}_\alpha(Y)$ is the quantile q at level α . The ES alone is not elicitable, meaning there is no scoring rule or loss function that can uniquely identify the ES as the optimal optimiser. However, Fissler et al. (2015) showed that the quantile function and ES are jointly elicitable, minimising the Fissler-Ziegel class of loss functions in expectation. One of which takes the form:

$$L_\alpha^{\text{FZ}}(y, q, e) = (\mathbb{I}(y \leq q) - \alpha)(q - y) + \frac{1}{\alpha} \exp(e) \mathbb{I}(y \leq q)(q - y) + \exp(e)(e - q - 1),$$

where y is the realisation of Y , q is the quantile, and e is the ES. This loss function is non-convex, thus greatly complicating the optimisation problem. Alternatively, we propose to use the Rockafellar-Uryasev joint characterisation of the quantile and ES (Rockafellar & Uryasev, 2000): $\tilde{\text{ES}}_k(\mathbf{P}) = \max_q E_{\tilde{\mathbf{y}}} [H_\alpha(\tilde{\mathbf{y}}_k(\mathbf{P}), q)]$ and $\tilde{q}_k(\mathbf{P}) = \arg \max_q E_{\tilde{\mathbf{y}}} [H_\alpha(\tilde{\mathbf{y}}_k(\mathbf{P}), q)]$, where $H_\alpha(y, q) = q - \frac{1}{\alpha} \max\{0, y - q\}$. Thus, we jointly reconcile the quantile and ES using the following bilevel optimisation problem:

$$\begin{aligned}
& \min_{\mathbf{P}} \sum_{k=1}^K E_{\mathbf{y}} \left[L_{\alpha}^{\text{FZ}}(y_k, \tilde{q}_k(\mathbf{P}), \widetilde{\text{ES}}_k(\mathbf{P})) \right] \\
& \text{s.t. } \tilde{q}_k(\mathbf{P}) = \arg \max_q E_{\tilde{\mathbf{y}}} [H_{\alpha}(\tilde{y}_k(\mathbf{P}), q)], \\
& \quad \widetilde{\text{ES}}_k(\mathbf{P}) = \max_q E_{\tilde{\mathbf{y}}} [H_{\alpha}(\tilde{y}_k(\mathbf{P}), q)], \quad \forall k = 1, \dots, K.
\end{aligned} \tag{1}$$

This ensures that the lower-level problem is still convex, however the upper-level objective function with the Fissler-Ziegler loss is non-convex. Irrespective of this non-convexity, smoothing techniques combined with the implicit function and envelope theorems can be used to devise gradient-based methods for locally minimising.

4 Applications

5 Contributions

References

- Angam, B., Beretta, A., De Poorter, E., Duvinage, M., & Peralta, D. (2025). Forecast reconciliation for vaccine supply chain optimization. In *Communications in computer and information science* (pp. 101–118). Springer Nature Switzerland. https://doi.org/10.1007/978-3-031-74650-5/_6
- Athanasopoulos, G., Ahmed, R. A., & Hyndman, R. J. (2009). Hierarchical forecasts for australian domestic tourism. *International Journal of Forecasting*, 25(1), 146–166. <https://doi.org/10.1016/j.ijforecast.2008.07.004>
- Athanasopoulos, G., Hyndman, R. J., Kourentzes, N., & Panagiotelis, A. (2024). *Forecast reconciliation: A review*. 40(2), 430–456. <https://www.sciencedirect.com/science/article/pii/S0169207023001097>
- Basel Committee on Banking Supervision. (2019). *Explanatory note on the minimum capital requirements for market risk*. https://www.bis.org/bcbs/publ/d457_note.pdf
- Di Modica, C., Pinson, P., & Ben Taieb, S. (2021). Online forecast reconciliation in wind power prediction. *Electric Power Systems Research*, 190(106637), 106637. <https://doi.org/10.1016/j.epsr.2020.106637>
- El Gemayel, J., Lafarguette, R., Itd, K. M., et al. (2022). *United arab emirates: Technical assistance reportliquidity management and forecasting*.
- Fissler, T., Ziegel, J. F., & Gneiting, T. (2015). Expected shortfall is jointly elicitable with value at risk - implications for backtesting. *arXiv [q-Fin.RM]*. <http://arxiv.org/abs/1507.00244>
- Hyndman, R. J., Ahmed, R. A., Athanasopoulos, G., & Shang, H. L. (2011). Optimal combination forecasts for hierarchical time series. *Computational Statistics & Data Analysis*, 55(9), 2579–2589. <https://doi.org/10.1016/j.csda.2011.03.006>
- Hyndman, R. J., Lee, A. J., & Wang, E. (2016). Fast computation of reconciled forecasts for hierarchical and grouped time series. *Computational Statistics & Data Analysis*, 97, 16–32. <https://doi.org/10.1016/j.csda.2015.11.007>
- Li, H., Li, H., Lu, Y., & Panagiotelis, A. (2019). A forecast reconciliation approach to cause-of-death mortality modeling. *Insurance, Mathematics & Economics*, 86, 122–133. <https://doi.org/10.1016/j.insmatheco.2019.02.011>
- Panagiotelis, A., Gamakumara, P., Athanasopoulos, G., & Hyndman, R. J. (2023). Probabilistic forecast reconciliation: Properties, evaluation and score optimisation. *European Journal of Operational Research*, 306(2), 693–706. <https://doi.org/10.1016/j.ejor.2022.07.040>
- Rockafellar, R. T., & Uryasev, S. (2000). Optimization of conditional value-at-risk. *Journal of Risk*, 2, 21–42. <https://www.risknet.de/fileadmin/eLibrary/Rockafellar-Conditional-Value-at-Risk.pdf>
- Seaman, B., & Bowman, J. (2022). Applicability of the M5 to forecasting at walmart. *International Journal of Forecasting*, 38(4), 1468–1472. <https://doi.org/10.1016/j.ijforecast.2021.06.002>
- Wickramasuriya, S. L., Athanasopoulos, G., & Hyndman, R. J. (2019). Optimal forecast reconciliation for hierarchical and grouped time series through trace minimization. *Jour-*

nal of the American Statistical Association, 114(526), 804–819. <https://doi.org/10.1080/01621459.2018.1448825>

Yamai, Y., & Yoshida, T. (2005). Value-at-risk versus expected shortfall: A practical perspective. *Journal of Banking & Finance*, 29(4), 997–1015. <https://doi.org/10.1016/j.jbankfin.2004.08.010>